

Buy it once

A storefront with an agent that watches how you shop and argues for the thing that will actually last. It cites who owns the brand today, whether parts will exist in 2032, and what the thing costs per year of use rather than at the till.

You cannot tell from the product page

“The market is splintered. Everyone else is in a market for lemons.”

HACKER NEWS, LAPTOP QUALITY THREAD, 113 POINTS

R/BUYITFORLIFE

“Are there any good brands left not owned by private equity?”

Top of the month, repeatedly.

A FOUR-WEEK EXPERIMENT

One person testing four T-shirts, €20 to €150, by hand

Because the dataset does not exist.

INVISIBLE AT THE TILL

Bosch 94% charitable trust • Birkenstock bought out 2021

On no product page anywhere.

Four fields nobody else surfaces

COST PER YEAR

₹500

Price ÷ expected service life.
The honest unit of comparison.

OWNERSHIP TODAY



Family, trust, employee owned,
public, private equity. Graded,
never a rage filter.

TRAJECTORY

“...”

Only ever a cited quote. Never
Hardy's own assertion.

SERVICEABILITY

8.0

Repairability, parts horizon, what
the warranty really covers.

Where there is no record, the page says unknown instead of guessing.

Behaviour in, a case out

01

Browse

Views, searches, dwell.
Batched and non-
blocking.

02

Threshold

12 meaningful events, or
a real shift in interest.

03

Reason

Infer what this person is
actually solving.

04

Retrieve

Hybrid search over the
real catalog.

05

Grade

Is each candidate
actually supported?

06

Argue

A case citing the user's
own actions.

It ends in **buy this one**. Hardy never says don't buy. It redirects you to the thing worth buying.

A graph, not a prompt chain

FRAMEWORK

LangGraph, six nodes, two conditional edges, a bounded refine loop `src/agent/graph.py`

STOPS EARLY

If the signal is too thin it ends before retrieval, so there is **no spend on nothing**.

GRADES EVIDENCE

Weak candidates send it back to refine the query. Maximum twice.

GROUNDED

Retrieval is hard-filtered to the catalog. Generation cannot invent a product.

EVERY AI CALL

Routes through one file, `src/integrations/mesh.py`

Be smart about when you call the model

FIRES WHEN

12 meaningful events since the last recommendation, the top category shifting, the 08:00 IST digest, or an admin asking

SUPPRESSED WHEN

Fewer than 5 events, one ran in the last 10 minutes, the profile hash is unchanged, which is **the cache hit**

MODEL CALLS PER EVENT

0.700

Live on /debug, computed from the same tables the storefront uses.

CACHE KEY COVERS

4

Categories, products, queries, catalog version.
An edit invalidates it.

DIGEST OVER STALE BEHAVIOUR

0

LLM calls. There is a test asserting it.

Two stores that cannot drift

ON WRITE

Create, edit and delete each write SQLite and Qdrant together. If the embedding fails, neither is kept.

PROOF, NOT A CLAIM

`GET /api/admin/consistency` counts both stores and names anything missing or orphaned.

LIVE RIGHT NOW

`136 = 136, in_sync true, 0 missing, 0 orphaned`

REPAIRABLE

One button re-embeds the whole catalog and rebuilds the vector store.

All four, implemented and running

★ STRUCTURED AGENT

LangGraph

Six nodes, conditional edges, and a bounded refine loop.

★ PROACTIVE DELIVERY

APScheduler

A real cron digest at 08:00 IST, not a manual button.

★ OBSERVABILITY

LangSmith

The whole graph traced end to end.

★ RETRIEVAL POLISH

Re-rank

Metadata filter plus a durability re rank, shown before and after.

Most entries land one or two. The brief calls these the difference between solid and exceptional.

Boring where it counts

BACKEND

FastAPI · Python 3.11 · domain-driven layout

FRONTEND

Jinja2 templates and vanilla JS for tracking. One process, no CORS

STORES

SQLite via SQLAlchemy. Qdrant at 1536 dimensions, cosine

LLM GATEWAY

Mesh API, running `openai/gpt-4o` and `text-embedding-3-small`

RUNNING ON

One 1GB droplet, three containers behind Caddy, TLS from Let's Encrypt

Same skeleton, opposite instinct

Every other entry builds an agent that persuades you to buy **more**. Hardy uses the identical mandated architecture: behavioural tracking, dual write to SQL and a vector database, agentic retrieval, persuasive generated copy, every call through Mesh. It points that at the **buyer** instead of the seller.

NOT

An affiliate site. No commissions, no sponsored placement.

NOT

Anti-consumption. It never says don't buy.

NOT

A review site. Ownership, serviceability, cost per year, not stars.

hardy.shashidev.me

REPOSITORY

github.com/shashiKundur1/hardy

Every line evaluated is in it.

GLASS BOX

/debug

Call counters, cache hit ratio, and why the last trigger fired.

CHECKS

4/4 critical

Compiles, dependencies, Mesh used, Mesh key valid.